

Home / Dashboard / My Courses / Microsoft Azure Exam AZ-204 Certification / Practice Test III / Report

 Microsoft Azure Exam AZ-204 Certification

Level: Advanced Back to the Course

Practice Test III Completed on Mon, 23 May 2022

 1st Attempt	 32/55 Marks Obtained	 58.18% Your Score
 0h 54m 16s Time Taken	 FAIL Result	

Domain wise Quiz Performance Report

No.	Domain	Total Question	Correct	Incorrect	Unattempted	Marked for Review
1	Develop Azure compute solutions	13	7	6	0	0
2	Develop for Azure storage	18	12	6	0	1
3	Implement Azure security	1	1	0	0	0
4	Monitor, troubleshoot, and optimize Azure solutions	9	4	5	0	0
5	Connect to and consume Azure services and third-party services	14	8	6	0	1
Total	All Domains	55	32	23	0	2

Join us on Slack community

Review the Answers Filter By Incorrect (23)

Question 1 Incorrect

Domain: Connect to and consume Azure services and third-party services

A software company is developing a software solution. The software solution is for a food delivery-based company. The software needs to adhere to the following workflow

- A driver selects the restaurants for which they will deliver orders.
- Orders are sent to all available drivers in an area.
- Only orders for the selected restaurants will appear for the driver.
- The first driver to accept an order removes it from the list of available orders.

The application needs to make use of the Azure Service Bus service.

Which of the following actions would you implement for this requirement? Choose 3 answers from the options given below

A. Create a Service Bus topic for each restaurant for which a driver can receive messages. wrong

B. Create a single Service Bus topic right

C. Create a single Service Bus subscription

D. Create a single Service Bus Namespace right

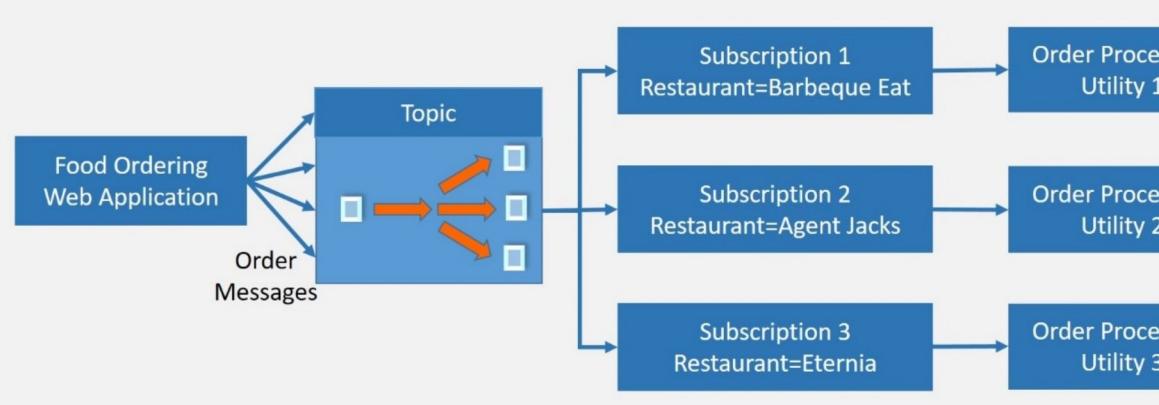
E. Create a Service Bus Namespace for each restaurant for which a driver can receive messages.

F. Create a Service Bus Subscription for each restaurant for which a driver can receive messages. right

Explanation:

Answer – B, D and F

We need to create Namespace first then Topics and finally Subscription as following flow chart



You should first create a Service Bus Namespace. Option E is incorrect since creating a Namespace for each restaurant would just be a maintenance overhead and difficult to keep track of via a program.

Here since the driver needs to choose the restaurant, that means the driver can be a subscriber.

Here you should have just one Topic. If you have multiple topics, then an order needs to be sent to all topics. Then deleting an order once it has been picked by a driver will be an issue. So, Option A gets ruled out.

You can create subscriptions and create rules based on driver and area.

For more information on Azure Service Bus, one can go to the below link

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-messaging-overview>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 5

Incorrect

Domain: Monitor, troubleshoot, and optimize Azure solutions

A development team has published an ASP.NET Web Application to the Azure Web App Service. They are also using Application Insights for the Web App for monitoring purposes and want to monitor telemetry data. They have to ensure that the cost of Application Insights does not exceed a pre-set budget. Which of the following would you implement to adhere to this requirement?

- A. Implement ingestion sampling using the Azure portal.
- B. Set a daily cap for the Application Insights instance.
- C. Implement adaptive sampling using the Azure portal.
- D. Implement adaptive sampling using the Application Insights SDK. right
- E. Implement ingestion sampling using the Application Insights SDK. wrong

Explanation:

Answer – D

Adaptive sampling is the default for the ASP.NET SDK. Adaptive sampling automatically adjusts to the volume of telemetry that your app sends. It operates automatically in the SDK in your web app so that telemetry traffic on the network is reduced.

Option B: Daily Cap is also a possible answer, but it has a major disadvantage. In case of data capacity is exhausted at any point during the day, then rest of the day, no data will be ingested and hence in case of any problem in the VM will not be notified. Hence this is not recommended.

For more information on managing costs for Application Insights, one can go to the below link

<https://docs.microsoft.com/en-us/azure/azure-monitor/app/pricing#managing-your-data-volume>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 6

Incorrect

Domain: Monitor, troubleshoot, and optimize Azure solutions

A development team is developing an application. The application will be working with customer data. The application will also be making use of Azure Redis Cache. You need to invalidate the cache when the customer data is changed.

You have to complete the below code to comply with the requirement

```
void clearCustomerCache(string p_Customer)
{
    //Establish the cache connection
    Slot1
    //Invalidate the cache
    Slot2
}
```

Which of the following will go into Slot1?

- A. IDatabase cache=Connection.GetDatabase(); right
- B. IDatabase cache=Connection.GetCache();
- C. ICache cache=Connection.GetDatabase();
- D. ICache cache=Connection.GetCache(); wrong

Explanation:

Answer – A

The right way is to use the IDatabase interface. Also you need to use the GetDatabase() method. This is also mentioned in the Microsoft documentation.

```
static void Main(string[] args)
{
    // Connection refers to a property that returns a ConnectionMultiplexer
    // as shown in the previous example.
    IDatabase cache = lazyConnection.Value.GetDatabase();
```



Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on an example on how to work with Azure Redis from .Net, one can go to the below link

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 8

Incorrect

Domain: Monitor, troubleshoot, and optimize Azure solutions

You have to load data into a cache in order to improve the performance and maintain data consistency between data store and data in cache. Which of the following are suitable scenarios to use Cache-Aside pattern? Choose 2 answers from the options given below.

- A. Static cached data which can be fit on startup wrong
- B. Native read-through and write-through cache operations are not provided right
- C. Unpredictable resource demand. right
- D. Caching session state information hosted in a web farm. wrong

Explanation:

Correct Answers: B and C

Option B is CORRECT because Cache-Aside pattern is suitable for cache which doesn't provide read-through and write-through operation in native.

Option C is CORRECT because using the Cache-Aside pattern enables applications to load data on demand. If the resource demand is unpredictable, it is suggested to use the Cache-Aside pattern.

Option A is incorrect because if the cached data set is static, the data can be cached during startup and apply a policy to prevent the data from expiring.

Option D is incorrect because developers should avoid introducing dependencies when caching session state information in a web farm.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/patterns/cache-aside>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 9

Incorrect Marked for review

Domain: Connect to and consume Azure services and third-party services

0

A company is developing a solution that allows smart devices to send information to a central location. The solution must receive and store messages until they can be processed. You are requested to create an Azure Service Bus Namespace.

Which Azure CLI command or PowerShell CmdLet should you run?

- A. New-AzServiceBusNamespace -ResourceGroupName "whizlabs-rg" -Location "EastUS" -Name "whizlabsnamespace" -SkuName "Standard" right
- B. az group create --name "whizlabs-rg" --location "Central US"
- C. New-AzResourceGroup -Name "whizlabs-rg" -Location "Central US"
- D. New-AzServiceBusNamespace -ResourceGroup "whizlabs-rg" -NamespaceName "whizlabs" -Location "WestUS" -SkuName "Standard" wrong

Explanation:

Answer – A

We need to create an Azure Service Bus Namespace

The full syntax in powershell is given at

<https://docs.microsoft.com/en-us/powershell/module/az.servicebus/new-azservicebusnamespace?view=azps-6.3.0>

```
New-AzServiceBusNamespace
[-ResourceGroupName] <String>
[-Location] <String>
[-Name] <String>
[-SkuName <String>]
[-SkuCapacity <Int32>]
[-Tag <Hashtable>]
[-DefaultProfile <AzureContextContainer>]
[-WhatIf]
[-Confirm]
<CommonParameters>
```

Hence Option A is correct answer

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 11

Incorrect

Domain: Connect to and consume Azure services and third-party services

A developer has setup a web application in Azure and also setup Azure CDN to route requests to the Web App. One of the requirements is to ensure that if users make requests based on passing an ID parameter, then those requests should always be served from a Point of Presence. An example of the URL is given below

<https://whizlabs.com/Customer.aspx?ID=1>

Which of the following mode should be set for the query string setting for the CDN service?

- A. Ignore query strings
- B. Default setting
- C. Bypass caching wrong
- D. Cache every unique URL right

Explanation:

Answer - D

Below are the different settings available for the CDN when it comes to caching of the query string.

Since we need to ensure that query strings are cached, we have to choose the option of 'Cache every unique URL'

- **Ignore query strings:** Default mode. In this mode, the CDN point-of-presence (POP) node passes the query strings from the requestor to the origin server on the first request and caches the asset. All subsequent requests for the asset that are served from the POP ignore the query strings until the cached asset expires.
- **Bypass caching for query strings:** In this mode, requests with query strings are not cached at the CDN POP node. The POP node retrieves the asset directly from the origin server and passes it to the requestor with each request.
- **Cache every unique URL:** In this mode, each request with a unique URL, including the query string, is treated as a unique asset with its own cache. For example, the response from the origin server for a request for example.ashx?q=test1 is cached at the POP node and returned for subsequent caches with the same query string. A request for example.ashx?q=test2 is cached as a separate asset with its own time-to-live setting.

Since this is the ideal approach, all other options are incorrect

For more information on working with query strings for CDN, one can go to the below link

<https://docs.microsoft.com/en-us/azure/cdn/cdn-query-string>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?

**Question 16**

Incorrect

Domain: Develop Azure compute solutions

Your company has a set of Azure storage accounts. These storage accounts store blob objects. You have to move the blobs from one container to another across storage accounts.

You have to use the right tool to perform the movement of data.

You decide to use the Azure CLI tool to implement this requirement

Would this fulfil the requirement?

- A. Yes right
- B. No wrong

Explanation:

Answer - A

Yes, you can also use the Azure CLI tool for this requirement

The below command can be used for this requirement

az storage blob copy start

[Edit](#)

Copies a blob asynchronously. Use `az storage blob show` to check the status of the blobs.

```
az storage blob copy start --destination-blob
  --destination-container
  [--account-key]
  [--account-name]
  [--auth-mode {key, login}]
  [--connection-string]
  [--destination-if-match]
  [--destination-if-modified-since]
  [--destination-if-none-match]
  [--destination-if-unmodified-since]
  [--destination-lease-id]
  [--metadata]
```

For more information on the Azure CLI command, please refer to the below link

<https://docs.microsoft.com/en-us/cli/azure/storage/blob/copy?view=azure-cli-latest>

Ask our Experts

[View Queries](#)

Did you like this Question?



Question 17

Incorrect

Domain: Develop Azure compute solutions

Your company has a set of Azure storage accounts. These storage accounts store blob objects. You have to move the blobs from one container to another across storage accounts.

You have to use the right tool to perform the movement of data.

You decide to use the PowerShell tool to implement this requirement

Would this fulfil the requirement?

A. Yes right

B. No wrong

Explanation:

Answer – A

Yes, you can also use the PowerShell tool for this requirement

The below command can be used for this requirement

```
Start-AzStorageBlobCopy
[-SrcBlob <String>
-SrcContainer <String>
-DestContainer <String>
[-DestBlob <String>]
[-PremiumPageBlobTier <PremiumPageBlobTier>]
[-StandardBlobTier <String>]
[-RehydratePriority <RehydratePriority>]
[-Context <IStorageContext>]
[-DestContext <IStorageContext>]
[-Force]
[-ServerTimeoutPerRequest <Int32>]
[-ClientTimeoutPerRequest <Int32>]
[-DefaultProfile <AzureContextContainer>]
[-ConcurrentTaskCount <Int32>]
[-WhatIf]
[-Confirm]
[<CommonParameters>]
```

For more information on the PowerShell command, please refer to the below link

<https://docs.microsoft.com/en-us/powershell/module/az.storage/start-azstorageblobcopy?view=azps-4.6.0>

Ask our Experts

[View Queries](#)

Did you like this Question?



Question 21

Incorrect

Domain: Monitor, troubleshoot, and optimize Azure solutions

A developer has been assigned a task to create code which would interact with an Azure Redis instance. Objects of the following class need to be uploaded to the Azure Redis Cache database

```
class WhizlabCustomer
{
    public string Id { get; set; }
    public string Name { get; set; }

    public WhizlabCustomer(string pID, string pName)
    {
        this.Id = pID;
        this.Name = pName;
    }
}
```

You need to complete the below code snippet

```
// Code to store the object in cache
WhizlabCustomer obj = new WhizlabCustomer("1", "David");
cache. Slot1 ("ID1", JsonConvert. Slot2 (obj));

// Retrieve the object from the cache
WhizlabCustomer objcache = JsonConvert. Slot3 <WhizlabCustomer>(cache. Slot4 ("ID1"));
```

Which of the following would go into Slot1?



- A. ObjectSet wrong
- B. StringSet right
- C. ClassSet
- D. Set

Explanation:

Answer – B

To Add an object to the cache database, we need to use the StringSet Method. An example of this is also given in the Microsoft documentation

Add the following Employee class definition to Program.cs:

```
C#
class Employee
{
    public string Id { get; set; }
    public string Name { get; set; }
    public int Age { get; set; }

    public Employee(string EmployeeId, string Name, int Age)
    {
        this.Id = EmployeeId;
        this.Name = Name;
        this.Age = Age;
    }
}
```

At the bottom of Main() procedure in Program.cs, and before the call to Dispose(), add the following lines of code to cache and retrieve a serialized .NET object:

```
C#
// Store .NET object to cache
Employee e007 = new Employee("e007", "Davide Columbo", 100);
Console.WriteLine("Cache response from storing Employee .NET object : " +
    cache.StringSet("e007", JsonConvert.SerializeObject(e007))); -----^

// Retrieve .NET object from cache
Employee e007FromCache = JsonConvert.DeserializeObject<Employee>(cache.StringGet("e007"));
Console.WriteLine("Deserialized Employee .NET object :\n");
Console.WriteLine("\tEmployee.Name : " + e007FromCache.Name);
Console.WriteLine("\tEmployee.Id : " + e007FromCache.Id);
Console.WriteLine("\tEmployee.Age : " + e007FromCache.Age + "\n");
```

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on how to work with Azure Redis Cache from a .Net program, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question? Like Dislike

Question 24

Incorrect

Domain: Monitor, troubleshoot, and optimize Azure solutions

A developer has been assigned a task to create code which would interact with an Azure Redis instance. Objects of the following class need to be uploaded to the Azure Redis Cache database

```
class WhizlabCustomer
{
    public string Id { get; set; }
    public string Name { get; set; }

    public WhizlabCustomer(string pID, string pName)
    {
        this.Id = pID;
        this.Name = pName;
    }
}
```

You need to complete the below code snippet

```
// Code to store the object in cache
WhizlabCustomer obj = new WhizlabCustomer("1", "David");

cache. Slot1 ("ID1", JsonConvert. Slot2 (obj)));

// Retrieve the object from the cache
WhizlabCustomer objcache = JsonConvert. Slot3 <WhizlabCustomer>(cache. Slot4 ("ID1"));
```

Which of the following would go into Slot4?

- A. ObjectGet wrong
- B. StringSet
- C. ClassSet
- D. StringGet right



Explanation:

Answer – D

We need to get an object from the cache, so we should use the StringGet method.

An example of this is also given in the Microsoft documentation

Add the following `Employee` class definition to `Program.cs`:

```
C#
class Employee
{
    public string Id { get; set; }
    public string Name { get; set; }
    public int Age { get; set; }

    public Employee(string EmployeeId, string Name, int Age)
    {
        this.Id = EmployeeId;
        this.Name = Name;
        this.Age = Age;
    }
}
```

At the bottom of `Main()` procedure in `Program.cs`, and before the call to `Dispose()`, add the following lines of code to cache and retrieve a serialized .NET object:

```
C#
// Store .NET object to cache
Employee e007 = new Employee("007", "Davide Columbo", 100);
Console.WriteLine("Cache response from storing Employee .NET object : " +
    cache.StringSet("e007", JsonConvert.SerializeObject(e007)));

// Retrieve .NET object from cache
Employee e007FromCache = JsonConvert.DeserializeObject<Employee>(cache.StringGet("e007"));
Console.WriteLine("Deserialized Employee .NET object :\n");
Console.WriteLine("\tEmployee.Name : " + e007FromCache.Name);
Console.WriteLine("\tEmployee.Id : " + e007FromCache.Id);
Console.WriteLine("\tEmployee.Age : " + e007FromCache.Age + "\n");
```



Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on how to work with Azure Redis Cache from a .Net program, please visit the below URL

<https://docs.microsoft.com/en-us/azure/azure-cache-for-redis/cache-dotnet-how-to-use-azure-redis-cache>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?  

Question 25

[Incorrect](#)

Domain: Develop for Azure storage

A company is developing a system which is going to be using Azure Cosmos DB at the underlying data store. Below are the requirements of the data store

- Ensure at least 99.99% availability and provide low latency
- Accept reservations event when localized network outages or other unforeseen failures occur.
- Process data in the same sequence as the writes being made
- Allow out of order data with a maximum of 5 second tolerance window

You have to provision a Cosmos DB account – SQL API. You already have a resource group in the South Central US region.

You have to complete the below Azure CLI commands for this purpose.

```
resourceGroup='whizlabs-rg'
accountname = 'whizlabsacc'
databasename = 'whizlabsdb'
collectionName='whizlabscollection'

consistencyLevel = ' Slot1 '
```

az cosmosdb create --name \$accountname
--resource-group \$resourceGroup
--max-interval 5
--default-consistency-level = \$consistencyLevel
Slot2
--locations regionName=southcentralus failoverPriority=0 isZoneRedundant=False
--locations regionName= NorthCentralUS failoverPriority=1 isZoneRedundant=True
Slot3

Which of the following would go into Slot1?

- A. Strong
- B. Eventual wrong
- C. ConsistentPrefix
- D. BoundedStaleness right

Explanation:

Answer – D



Since you can have an out of order read for a maximum of 5 seconds, this becomes our staleness window.

The Microsoft documentation mentions the following on the Bounded Staleness consistency level.

- **Bounded staleness:** The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "K" versions (i.e., "updates") of an item or by "T" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
 - The number of versions (K) of the item
 - The time interval (T) by which the reads might lag behind the writes

Bounded staleness offers total global order except within the "staleness window." The monotonic read guarantees exist within a region both inside and outside the staleness window. Strong consistency has the same semantics as the one offered by bounded staleness. The staleness window is equal to zero. Bounded staleness is also referred to as time-delayed linearizability. When a client performs read operations within a region that accepts writes, the guarantees provided by bounded staleness consistency are identical to those guarantees by the strong consistency.

Since this is clearly given in the Microsoft documentation, all other options are incorrect

For more information on consistency levels, please visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?

Question 26

Incorrect

Domain: Develop for Azure storage

A company is developing a system which is going to be using Azure Cosmos DB as the underlying data store. Below are the requirements of the data store

- Ensure at least 99.99% availability and provide network failures
- Accepts writes via the application even in the case of network outages or any unforeseen failures
- Process data in the same sequence as the writes being made
- Allow out of order data with a maximum of 5 second tolerance window

You have to provision a Cosmos DB account – SQL API. You already have a resource group in the South Central US region.

You have to complete the below Azure CLI commands for this purpose.

```
resourceGroup='whizlabs-rg'
accountname = 'whizlabsacc'
databasename = 'whizlabsdb'
collectionName='whizlabscollection'

consistencyLevel = ' Slot1 '
```

az cosmosdb create --name \$accountname
 --resource-group \$resourceGroup
 --max-interval 5
 --default-consistency-level = \$consistencyLevel
 Slot2
 --locations regionName=southcentralus failoverPriority=0 isZoneRedundant=False
 --locations regionName= NorthCentralUS failoverPriority=1 isZoneRedundant=True
 Slot3

Which of the following would go into Slot2?

- A. --enable-virtual-network true
- B. --enable-automatic-failover true right
- C. --kind 'GlobalDocumentDB' wrong
- D. --kind 'MongoDB'

Explanation:

Answer – B

Since we have to ensure that the data needs to be available even in the case of network outages or any unforeseen failures, we have to enable automatic failover.

The Microsoft documentation mentions the following

```
--enable-automatic-failover

Enables automatic failover of the write region in the rare event that the region is unavailable due to
an outage. Automatic failover will result in a new write region for the account and is chosen based
on the failover priorities configured for the account.
accepted values: false, true
```

Option A is incorrect since there is no mention in the question of requiring the database to be part of a virtual network

Option C is incorrect since the default API chosen for the database is the SQL API

Option D is incorrect since we need to create a Cosmos DB account with the SQL API



For more information on the Cosmos DB create command, please visit the below URL

<https://docs.microsoft.com/en-us/cli/azure/cosmosdb?view=azure-cli-latest#az-cosmosdb-create>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 27

Incorrect

Domain: Develop for Azure storage

A company is developing a system which is going to be using Azure Cosmos DB as the underlying data store. Below are the requirements of the data store

- Ensure at least 99.99% availability and provide network failures
- Accepts writes via the application even in the case of network outages or any unforeseen failures
- Process data in the same sequence as the writes being made
- Allow out of order data with a maximum of 5 second tolerance window

You have to provision a Cosmos DB account – SQL API. You already have a resource group in the South Central US region.

You have to complete the below Azure CLI commands for this purpose.

```
resourceGroup='whizlabs-rg'
accountname='whizlabacc'
databasename='whizlabdb'
collectionName='whizlabcollection'
consistencyLevel=' Slot1 '
az cosmosdb create --name $accountname \
Slot2
--resource-group $resourceGroup --max-interval 5 \
--default-consistency-level=$consistencyLevel
Slot3
```

Which of the following would go into Slot3?

- A. --locations 'southcentralus' wrong
- B. --locations 'eastus'
- C. --locations "SouthCentralUS=0", "NorthCentralUS=1" right
- D. --locations 'southcentralus=0'

Explanation:

Answer – C

Since we need to have additional regions for failover purposes, we need to add multiple locations to the Cosmos DB account.

And keep in mind for failover purpose, the region should be mapped correctly

The paired region for "South Central US" is "North Central US"

The other options are incorrect since they only have one location specified.

For more information on the Cosmos DB create command, please visit the below URL

<https://docs.microsoft.com/en-us/cli/azure/cosmosdb?view=azure-cli-latest#az-cosmosdb-create>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 31

Incorrect Marked for review

Domain: Develop for Azure storage

[View Case Study](#)

You need to formulate a query that would be used to get all items from container customer where the order price is \$9.99

```
select * from c IN Slot1 where Slot2 ="9.99"
```

Which of the following would go into Slot1?

- A. orders
- B. customer wrong
- C. course
- D. customer.orders right

Explanation:

Answer – D

We can use the IN clause to query data from JSON arrays

An example of an implementation of querying the data is given below

Data in the container

The screenshot shows the Azure Cosmos DB SQL API interface. On the left, the database 'whizlabdb' and container 'customer' are selected. The 'Items' section shows two documents:

```

1 select * from c
[ {
    "id": "1",
    "Name": "UserA",
    "orders": [
        {
            "course": "Big Data",
            "price": "9.99"
        }
    ],
    "ratings": {
        "3": "100",
        "4": "200",
        "5": "300"
    },
    "_rid": "Q30iAOTAj2gBAAAAAAA==",
    "_self": "dbs/Q30iAA==/colls/Q30iAOTAj2g=/docs/Q30iAOTAj2gBAAAAAAA==/",
    "_etag": "\"0a00d517-0000-0100-0000-5d35887c0000\"",
    "_attachments": "attachments/",
    "_ts": 1563789436
},
{
    "id": "2",
    "Name": "UserB",
    "orders": [
        {
            ...
        }
    ]
}
]

```

After executing the query

The screenshot shows the Azure Cosmos DB SQL API interface. The query 'select * from c IN customer.orders where c.price="9.99"' returns one document:

```

1 select * from c IN customer.orders where c.price="9.99"
[ {
    "course": "Big Data",
    "price": "9.99"
}
]

```

Since this is clear from the implementation, all other options are invalid

For more information on SQL queries for array objects, please visit the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/sql-query-object-array>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question? [Like](#) [Dislike](#)

Question 37

Incorrect

Domain: Develop Azure compute solutions

[View Case Study](#)

The Azure Web App needs to be running at all times. You have to suggest the most cost effective plan for the Azure Web App.

You recommended "Standard" as Web App Service Plan.

Would this fulfil the requirement?

A. Yes wrong

B. No right

Explanation:

Answer – B

To fulfil this requirement, the App service plan must support the "Always On" feature. And this is supported with the Standard App Service Plan. But this is not the most cost-effective App Service plan. Since the Basic App Service Plan already has this feature, you should choose that plan to cut on costs.

Below is the snippet of the features of the various App Service Plans

	FREE	SHARED	BASIC	STANDARD	PREMIUM	ISOLATED*	APP SERVICE LINUX	CONSUMPTION PLAN (FUNCTIONS)
- Settings								
64-bit			✓	✓	✓	✓	✓	✓
App Service Advisor *			✓	✓	✓	✓	✓	
Always On			✓	✓	✓	✓		

For more information on the App Service Plans, please visit the below URL

<https://azure.microsoft.com/en-us/pricing/details/app-service/plans/>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 39

Incorrect

Domain: Develop for Azure storage

[View Case Study](#)

A lease needs to be applied on common blob's in the Azure storage account.

How would you complete the below REST API call for implementing a blob lease?

<https://whizlabstore2020.blob.core.windows.net/demo/whizlabcommon.json?>

Slot1

Which of the following should go into Slot1?

- A. `getlease`
- B. `comp=lease` right
- C. `get=lease`
- D. `obj=lease` wrong

Explanation:

Answer - B

We have to use the query string parameter as "comp=lease"

This is also given in the Microsoft documentation

Request

The `Lease Blob` request may be constructed as follows. HTTPS is recommended. Replace `myaccount` with the name of your storage account:

PUT Method Request URI	HTTP Version
https://myaccount.blob.core.windows.net/mycontainer/myblob?comp=lease	HTTP/1.1

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with Blob lease, please visit the below URL

<https://docs.microsoft.com/en-us/rest/api/storageservices/lease-blob>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 41

Incorrect

Domain: Develop for Azure storage

A company is developing a series of applications. Each of these applications would be interacting with separate Azure CosmosDB accounts. Each application has a different requirement when it comes to accessing the underlying data. You have to set the consistency level for the Azure CosmosDB accounts based on each application requirement. You have to choose the most cost-effective consistency level for each CosmosDB account. Below are the requirements for each application when it comes to the consistency of the underlying data

Application Name	Data requirement
whizlabappA	This is not a critical application hence here no ordering of reads is required.
whizlabappB	Here the application users must never see out-of-order writes
whizlabappC	Here the users must always see the latest committed write.
whizlabappD	Here the data can be stale by at most 2 versions

Which of the following would you choose as the consistency level for the CosmosDB account used by the application "whizlabappB"?



- A. Strong
 B. Bounded Staleness
 C. Session wrong
 D. Consistent prefix right
 E. Eventual

Explanation:

Answer - D

Since here the requirement is that the user must never see out-of-order writes, the most cost-effective option is to use the "Consistent prefix" consistency level.

The Microsoft documentation mentions the following

Guarantees associated with consistency levels

The comprehensive SLAs provided by Azure Cosmos DB guarantee that 100 percent of read requests meet the consistency guarantee for any consistency level you choose. A read request meets the consistency SLA if all the consistency guarantees associated with the consistency level are satisfied. The precise definitions of the five consistency levels in Azure Cosmos DB using the TLA+ specification language are provided in the [azure-cosmos-tla](#) GitHub repo.

The semantics of the five consistency levels are described here:

- **Strong:** Strong consistency offers a linearizability guarantee. Linearizability refers to serving requests concurrently. The reads are guaranteed to return the most recent committed version of an item. A client never sees an uncommitted or partial write. Users are always guaranteed to read the latest committed write.
- **Bounded staleness:** The reads are guaranteed to honor the consistent-prefix guarantee. The reads might lag behind writes by at most "K" versions (i.e., "updates") of an item or by "T" time interval. In other words, when you choose bounded staleness, the "staleness" can be configured in two ways:
 - The number of versions (K) of the item
 - The time interval (T) by which the reads might lag behind the writes
- **Session:** Within a single client session reads are guaranteed to honor the consistent-prefix (assuming a single "writer" session), monotonic reads, monotonic writes, read-your-writes, and write-follows-reads guarantees. Clients outside of the session performing writes will see eventual consistency.
- **Consistent prefix:** Updates that are returned contain some prefix of all the updates, with no gaps. Consistent prefix consistency level guarantees that reads never see out-of-order writes.
- **Eventual:** There's no ordering guarantee for reads. In the absence of any further writes, the replicas eventually converge.

Since this is clearly given in the documentation, all other options are incorrect

For more information on CosmosDB consistency levels, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/cosmos-db/consistency-levels>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?  

Question 44

Incorrect

Domain: Develop Azure compute solutions

A company has a web application deployed to Azure. The web application is currently being hosted as part of the Azure Web App service. There is a requirement to stream the logs from the web app and filter out any errors.

You have to complete the below Azure CLI command for this requirement

```
az webapp Slot1 Slot2 --name "whizlabapp"
--resource-group "whizlab-rg" Slot3 Errors
```

Which of the following would go into Slot1?

- A. file
 B. log right
 C. tail
 D. stream wrong

Explanation:

Answer – B

Since we need to get the log files for the Azure Web app, we have to use the "log" option in the command.

An example of this is also given in the Microsoft documentation



Streaming with Azure CLI

To stream logging information, open a new command prompt, PowerShell, Bash, or Terminal session and enter the following command:

```
az webapp log tail --name appname --resource-group myResourceGroup
```

This command connects to the app named 'appname' and begin streaming information to the window as log events occur on the app. Any information written to files ending in .txt, .log, or .htm that are stored in the /LogFiles directory (d:/home/logfiles) is streamed to the local console.

To filter specific events, such as errors, use the **--Filter** parameter. For example:

```
az webapp log tail --name appname --resource-group myResourceGroup --filter Error
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with diagnostics logs, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs>



Question 45

Domain: Develop Azure compute solutions

A company has a web application deployed to Azure. The web application is currently being hosted as part of the Azure Web App service. There is a requirement to stream the logs from the web app and filter out any errors.

You have to complete the below Azure CLI command for this requirement

```
az webapp Slot1 Slot2 --name "whizlabapp"  
--resource-group "whizlab-rg" Slot3 Errors
```

Which of the following would go into Slot2?

- A. file
- B. log wrong
- C. tail right
- D. stream

Explanation:

Answer – C

Since we need to stream the log files, the next option to include is the "tail" option

An example of this is also given in the Microsoft documentation

Streaming with Azure CLI

To stream logging information, open a new command prompt, PowerShell, Bash, or Terminal session and enter the following command:

```
az webapp log tail --name appname --resource-group myResourceGroup
```

This command connects to the app named 'appname' and begin streaming information to the window as log events occur on the app. Any information written to files ending in .txt, .log, or .htm that are stored in the /LogFiles directory (d:/home/logfiles) is streamed to the local console.

To filter specific events, such as errors, use the **--Filter** parameter. For example:

```
az webapp log tail --name appname --resource-group myResourceGroup --filter Error
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with diagnostics logs, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs>



Question 46

Domain: Develop Azure compute solutions



A company has a web application deployed to Azure. The web application is currently being hosted as part of the Azure Web App service. There is a requirement to stream the logs from the web app and filter out any errors.

You have to complete the below Azure CLI command for this requirement

```
az webapp log tail --name "whizlabapp"
--resource-group "whizlab-rg" --slot3 Errors
```

Which of the following would go into Slot3?

- A. --path
- B. --only-show-right
- C. --key
- D. --type wrong

Explanation:

Answer – B

Since we have the filter out on the errors, we can use the filter option for the command.

An example of this is also given in the Microsoft documentation

```
az webapp log tail [--ids]
  [--name]
  [--provider]
  [--resource-group]
  [--slot]
  [--subscription]

Optional Parameters
--ids
One or more resource IDs (space-delimited). It should be a complete resource ID containing all information of 'Resource Id' arguments. You should provide either --ids or other 'Resource Id' arguments.

--name -n
Name of the web app. If left unspecified, a name will be randomly generated. You can configure the default using az configure --defaults web=<name>.

--provider
By default all live traces configured by az webapp log config will be shown, but you can scope to certain providers/folders, e.g. 'application', 'http', etc. For details, check out https://github.com/projectkudu/kudu/wiki/Diagnostic-Log-Stream.

--resource-group -g
Name of resource group. You can configure the default group using az configure --defaults group=<name>.

--slot -s
The name of the slot. Default to the production slot if not specified.

--subscription
Name or ID of subscription. You can configure the default subscription using az account set -s NAME_OR_ID.

Global Parameters
--debug
Increase logging verbosity to show all debug logs.

--help -h
Show this help message and exit.

--only-show-errors
Only show errors, suppressing warnings.

--output -o
Output format.

--query
JMESPath query string. See http://jmespath.org/ for more information and examples.

--verbose
Increase logging verbosity. Use --debug for full debug logs.
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on working with diagnostics logs, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/app-service/troubleshoot-diagnostic-logs>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?

Question 48

Incorrect

Domain: Connect to and consume Azure services and third-party services

A team has to integrate various modules of an application with the Azure Event Grid service. They have to filter events which are sent to the various application endpoints. The requirements for the type of messages that need to be received by the different endpoints are given below

Application Endpoint	Message requirement
EndpointA	Receives failure messages for any resources deployed to the Azure subscription



EndpointB	Receives messages whenever objects are added to a specific container in Azure Blob storage
EndpointC	Receive messages whenever data fields in the message has the value of "Organization"

Which of the following would you use as a filter option for messages that need to be sent to EndpointB?

- A. Subject begins with or ends with right
- B. Advanced fields and operators
- C. ResourceTypes wrong
- D. EventTypes

Explanation:

Answer - A

Since here we need to check on the messages sent to a container, so we have to check the subject of the message

The Microsoft documentation mentions the following

Subject filtering

For simple filtering by subject, specify a starting or ending value for the subject. For example, you can specify the subject ends with .txt to only get events related to uploading a text file to storage account. Or, you can filter the subject begins with /blobServices/default/containers/testcontainer to get all events for that container but not other containers in the storage account.

When publishing events to custom topics, create subjects for your events that make it easy for subscribers to know whether they're interested in the event. Subscribers use the subject property to filter and route events. Consider adding the path for where the event happened, so subscribers can filter by segments of that path. The path enables subscribers to narrowly or broadly filter events. If you provide a three segment path like /A/B/C in the subject, subscribers can filter by the first segment /A to get a broad set of events. Those subscribers get events with subjects like /A/B/C or /A/D/E. Other subscribers can filter by /A/B to get a narrower set of events.

The JSON syntax for filtering by subject is:

```
JSON Copy
{
  "filter": {
    "subjectBeginsWith": "/blobServices/default/containers/mycontainer/log",
    "subjectEndsWith": ".jpg"
  }
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on event filtering in Azure Event Grid, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/event-grid/event-filtering>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question? like dislike

Question 49

Incorrect

Domain: Connect to and consume Azure services and third-party services

A team has to integrate various modules of an application with the Azure Event Grid service. They have to filter events which are sent to the various application endpoints. The requirements for the type of messages that need to be received by the different endpoints are given below

Application Endpoint	Message requirement
EndpointA	Receives failure messages for any resources deployed to the Azure subscription
EndpointB	Receives messages whenever objects are added to a specific container in Azure Blob storage
EndpointC	Receive messages whenever data fields in the message has the value of "Organization"

Which of the following would you use as a filter option for messages that need to be sent to EndpointC?

- A. Subject begins with or ends with wrong
- B. Advanced fields and operators right
- C. ResourceTypes
- D. EventTypes

Explanation:

Answer - B



Here since we need a more advanced scenario and check for the data field values, we have to choose the "Advanced fields and operators" filter option.

The Microsoft documentation mentions the following

Advanced filtering

To filter by values in the data fields and specify the comparison operator, use the advanced filtering option. In advanced filtering, you specify the:

- operator type - The type of comparison.
- key - The field in the event data that you're using for filtering. It can be a number, boolean, or string.
- value or values - The value or values to compare to the key.

The JSON syntax for using advanced filters is:

```
JSON Copy
{
  "filter": {
    "advancedFilters": [
      {
        "operatorType": "NumberGreaterThanOrEqualTo",
        "key": "Data.Key1",
        "value": 5
      },
      {
        "operatorType": "StringContains",
        "key": "Subject",
        "values": ["container1", "container2"]
      }
    ]
  }
}
```

Since this is clearly given in the documentation, all other options are incorrect

For more information on event filtering in Azure Event Grid, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/event-grid/event-filtering>

[Ask our Experts](#)

[View Queries](#)

Did you like this Question?



Question 53

Incorrect

Domain: Connect to and consume Azure services and third-party services

A development team needs to develop an application module that needs to interact with an Azure service bus queue.

Below is the snippet of code that needs to be completed. The code is used to send a message to the queue.

```
const string whizlabServiceBusConnectionString = "...";
const string QueueName = "whizlabqueue";
static IQueueClient whizlabqueueClient;
static async Task SideSync()
{
    whizlabqueueClient = new Slot1 (ServiceBusConnectionString, QueueName);
    await SendMessagesAsync(numberOfMessages);
    await whizlabqueueClient.CloseAsync();
}

static async Task SendMessagesAsync(string messageBody)
{
    try
    {
        var message = new Message(Encoding.UTF8.GetBytes(messageBody));
        await whizlabqueueClient. Slot2 (message);
    }
    catch (Exception whizlabexception)
    {
        Console.WriteLine($"{DateTime.Now} :: Exception:{whizlabexception.Message}");
    }
}
```

Which of the following needs to go into Slot1?

- A. Client
- B. ServiceBusClient wrong
- C. QueueClient right

D. BusClient

Explanation:

Answer – C

Since we need to interact with a queue, we have to use the "QueueClient" class.

An example of this is given in the Microsoft documentation

4. Directly after `Main()`, add the following asynchronous `MainAsync()` method that calls the send messages method:

```
C#  
static async Task MainAsync()  
{  
    const int numberOfMessages = 10;  
    QueueClient queueClient = new QueueClient(ServiceBusConnectionString, QueueName);  
  
    Console.WriteLine("*****");  
    Console.WriteLine("Press ENTER key to exit after sending all the messages.");  
    Console.WriteLine("*****");  
  
    // Send messages.  
    await SendMessagesAsync(numberOfMessages);  
  
    Console.ReadKey();  
  
    await queueClient.CloseAsync();  
}
```



Since this is clearly given in the documentation, all other options are incorrect

For more information on working with queues in .Net, one can go to the below URL

<https://docs.microsoft.com/en-us/azure/service-bus-messaging/service-bus-dotnet-get-started-with-queues>

<https://github.com/jignesh24/AzureExamples/blob/master/AzureServiceBus-Queue/BasicExample/BasicReceiverQueueExample/Program.cs>

[Ask our Experts](#)[View Queries](#)

Did you like this Question?

[Finish Review](#)

Certification

Cloud Certification
Java Certification
PM Certification
Big Data Certification

Company

Become Our Instructor
Support
Discussions
Blog
Business

Support

Contact Us
Help Topics

Join us on Slack!

Join our open [Slack community](#) and get your queries answered instantly! Our experts are online to answer your questions!

